



# Monitoring





**Monitoring** gives insights into your applications that help you detect, investigate, and remediate problems faster.





- **Monitoring** tracks resource health and usage through data collection.
- It provides **real-time insights** to identify operational issues.
- **Metrics** help address performance, availability, and capacity concerns.





s3

ec2



- **Amazon S3 Metrics:**

- Size of objects in a bucket.
- Number of objects in a bucket.
- Number of HTTP requests to a bucket.

- **Amazon RDS Metrics:**

- Database connections.
- CPU utilization.
- Disk space usage.

- **Amazon EC2 Metrics:**

- CPU utilization.
- Network utilization.
- Disk performance.
- Status checks.



SEEDMETE



RESOURCE



PROACTIVE ISSUE



COST BANK

## Benefits

- Detect and resolve issues proactively.
- Enhance resource performance and reliability.
- Identify and address security threats.
- Optimize operations with data-driven insights.
- Reduce costs through effective monitoring.



## Amazon CloudWatch

Amazon CloudWatch is a monitoring and observability service that collects your resource data and provides actionable insights into your applications.





## **You can use CloudWatch to do the following:**

- Detect anomalous behavior in your environments.
- Set alarms to alert you when something is not right.
- Visualize logs and metrics with the AWS Management Console.
- Take automated actions like scaling.
- Troubleshoot issues.
- Discover insights to keep your applications healthy.



## How CloudWatch Works

- **Managed Service:** Centralized monitoring without managing infrastructure.
- **Basic Monitoring:** Free, 1 data point per 5 minutes; suitable for most apps.
- **Detailed Monitoring:** 1-minute granularity at additional cost.
- **Metrics:** Time-ordered data points representing variables (e.g., CPU usage).
- **Dimensions:** Name-value pairs for filtering metrics (e.g., InstanceId).
- **Custom Metrics:** Record application-level data (e.g., page views, error rates).
- **High-Resolution Metrics:** 1-second granularity for precise tracking.







## CloudWatch Dashboards

- **Customizable Views:** Visualize multiple metrics in widgets (e.g., graphs, text).
- **Global Insights:** Combine metrics from different AWS Regions.
- **Live Data:** Display recent data published within the last minute.
- **External Tools:** Use APIs to ingest metrics into third-party tools.
- **Access Control:** Use IAM policies to manage dashboard permissions.



# Amazon CloudWatch Logs

- **Centralized Storage:** Store and analyze logs from AWS and on-premises apps.
- **Query & Filter Logs:** Search for specific data (e.g., error stack traces).
- **Metric Filters:** Convert log data into metrics for monitoring.
- **Log Agents:** Send data from EC2 instances using the CloudWatch Logs agent.
- **Source Examples:** Lambda logs (minimal setup), EC2 logs (requires agent).



## CloudWatch Logs Terminology

- Log Events: Individual data entries in log files.
- Log Streams: Sequence of log events from a specific source (e.g., EC2 instance).
- Log Groups: Collection of log streams sharing the same retention policies.



## CloudWatch Alarms

- **Triggering Actions:** Automatically respond to metric threshold breaches.
- States:
- **OK:** Metric within threshold; normal operation.
- **ALARM:** Metric exceeds threshold; potential issue.
- **INSUFFICIENT\_DATA:** Insufficient data to determine state.







## Steps to Create Alarms:

- Set up a metric filter (e.g., HTTP 500 errors).
- Define an alarm (e.g., threshold for 500 errors over 5 minutes).
- Define an action (e.g., send email or auto-remediate).
- Integrated Responses: Trigger EC2 actions, scaling, or Lambda functions.





## Prevent & Troubleshoot Issues

- Notification Example: Alert for HTTP 500 errors via Amazon SNS.
- Automated Actions: Auto-reboot EC2, scale resources, or invoke Lambda functions.
- Faster Remediation: Use alarms to preemptively address operational problems.





## Solution Optimization

Design your system to have no single point of failure by using automated monitoring, failure detection, and failover mechanisms.



# Understanding Availability

## Availability Metrics:

- Expressed in **percentages** or "nines" (e.g., 99.9% = 8.77 hrs/year downtime).

## Why It Matters:

- High availability ensures uninterrupted customer access to applications.

## Current Issue:

- Single EC2 instance = Single point of failure.

## Solution:

- Deploy a second EC2 in a different Availability Zone for redundancy.

## Challenges:

- Replication, customer redirection, and choosing availability type (active-passive vs. active-active).

# High Availability with Multiple Servers

## **Benefits:**

- Reduces risk from hardware or data center failures.

## **Customer Redirection:**

- Use DNS or Load Balancer for smooth traffic routing.

## **Active-Passive Systems:**

- Only one server active; suitable for stateful apps.

## **Active-Active Systems:**

- Both servers active; ideal for stateless apps.

## **Cost vs. Availability:**

- Higher availability means more infrastructure costs.

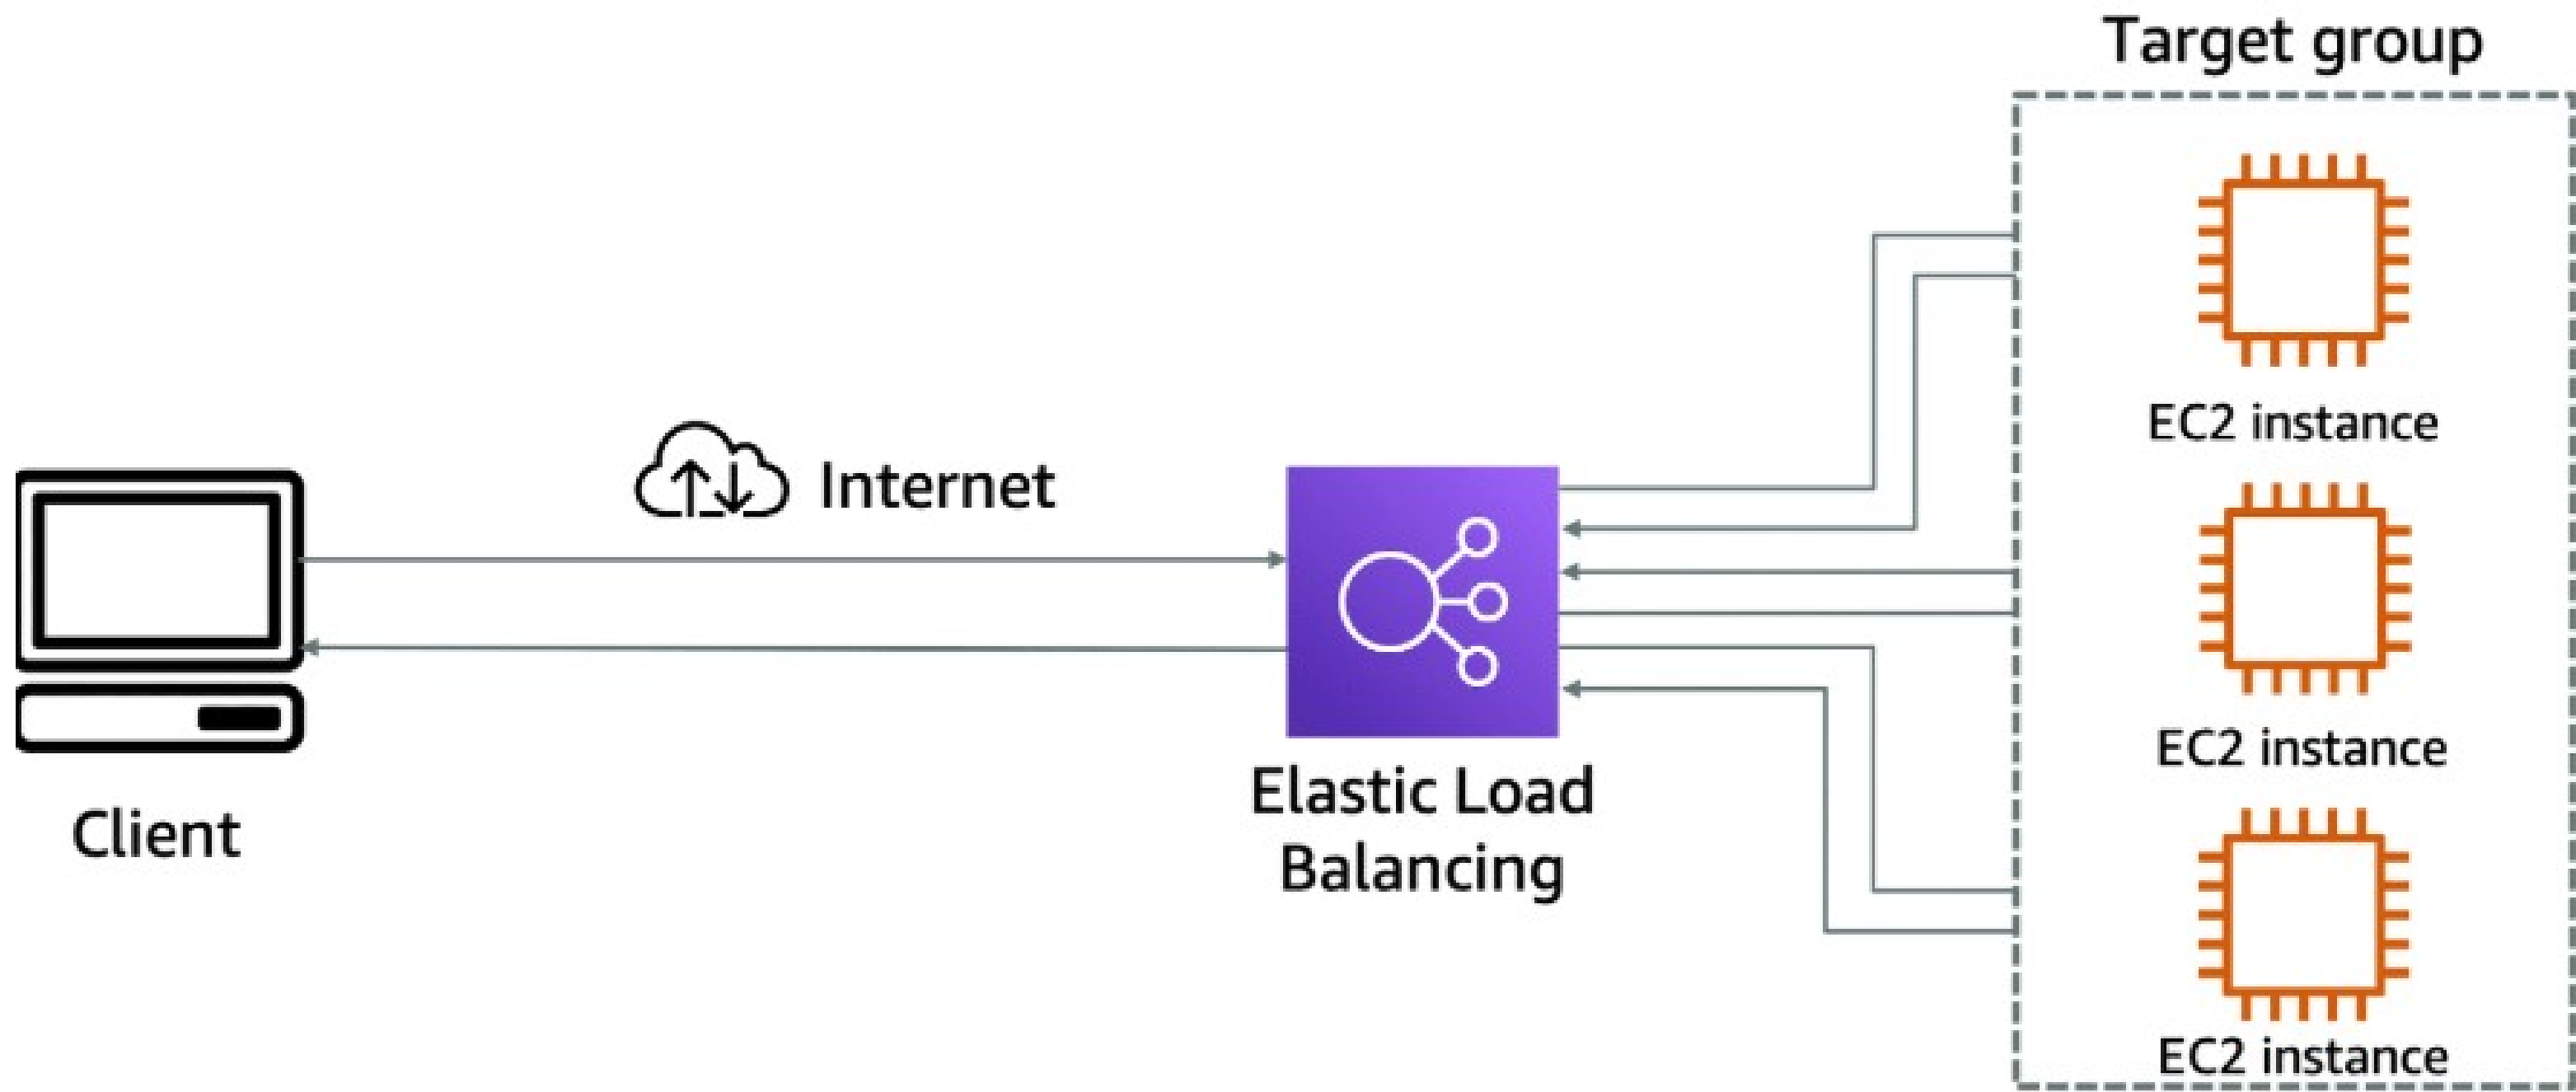


## Traffic Routing with Elastic Load Balancing

The Elastic Load Balancing (ELB) service can distribute incoming application traffic across EC2 instances, containers, IP addresses, and Lambda functions.











# Health Checks in Load Balancers (ELB)

## Health Check Types:

- TCP: Verifies connection to EC2 instance.
- HTTP/HTTPS: Checks response from specified URLs (e.g., /monitor).

## Purpose:

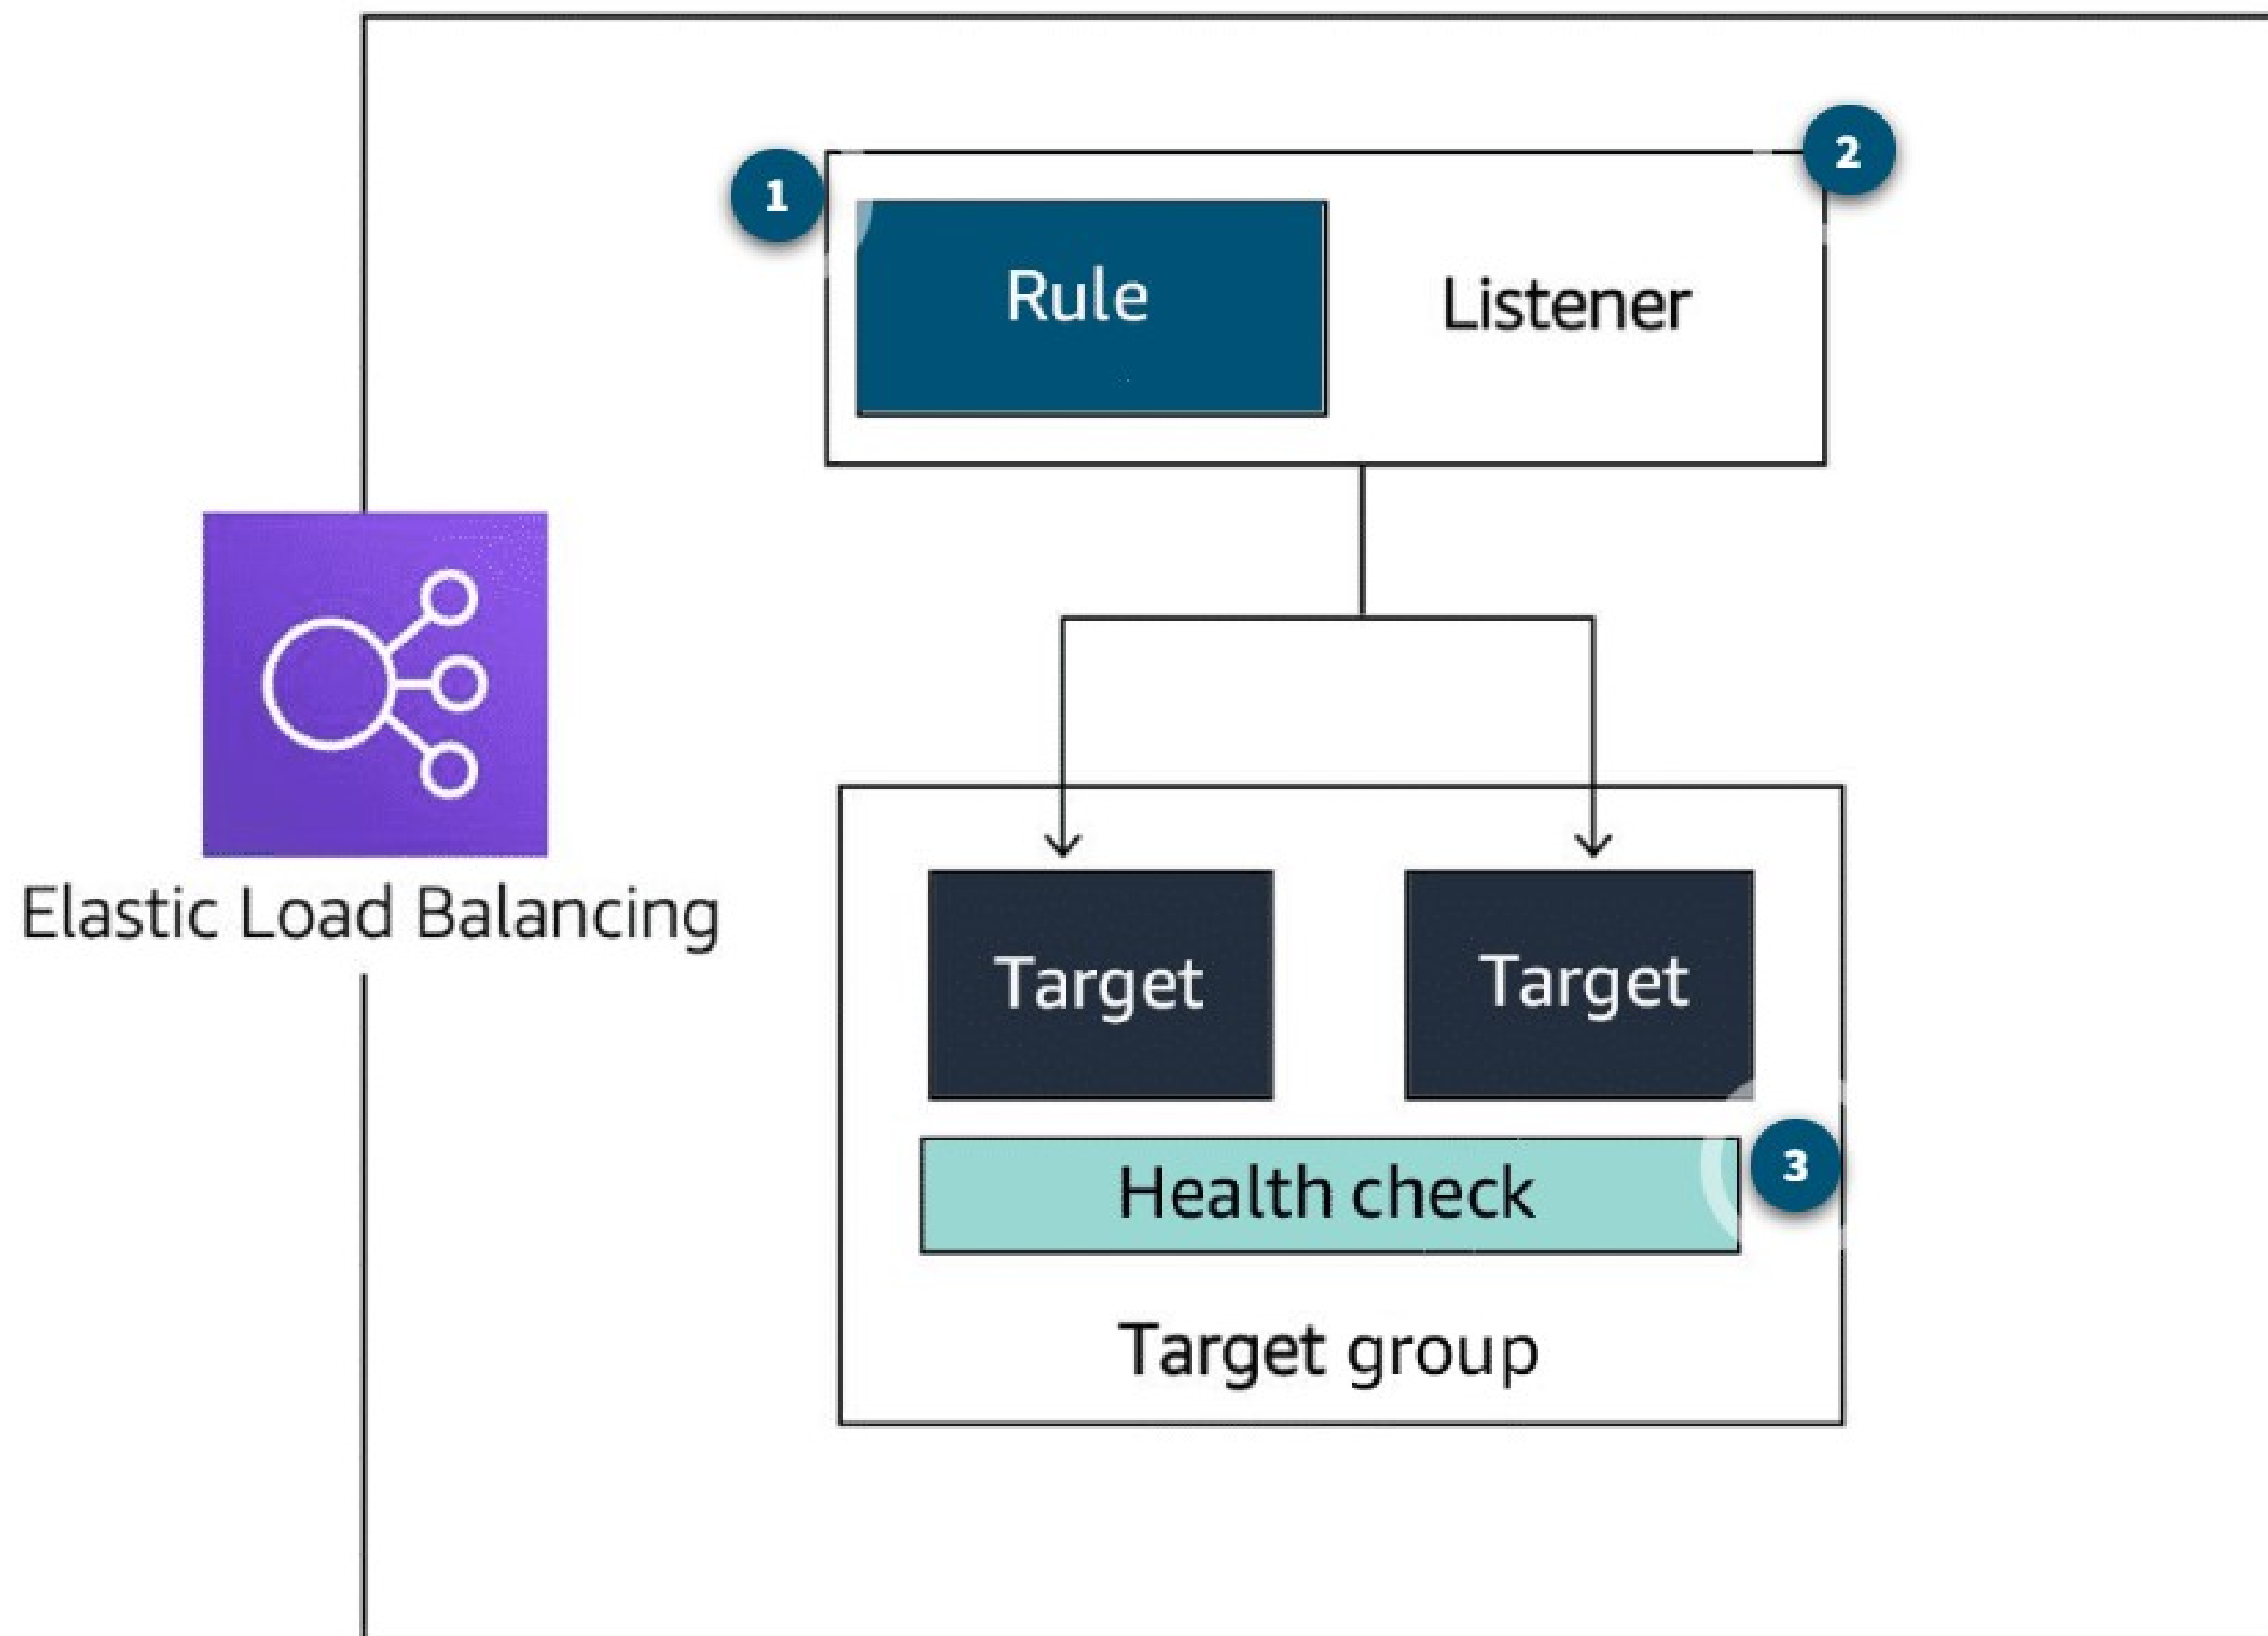
- Routes traffic only to healthy EC2 instances.

## Auto Scaling Integration:


- ELB informs Auto Scaling to replace unhealthy instances.

## Connection Draining:

- Prevents termination of instances with active connections.







### Application Load Balancer (ALB)

- User authorization
- Rich metrics & logging
- Redirects
- Fixed response

### Network Load Balancer (NLB)


- TCP & UDP connection-based
- Source IP preservation
- Low latency

### Gateway Load Balancer (GLB)

- Health checks
- Gateway Load Balancer Endpoints
- Higher availability for third-party virtual appliances

## AUTO-SCALING

### Automatic-Scaling




Dynamic – real-time adjustments  
Automated resource allocation

V S.

## TRADITIONAL-SCALING

### Traditional Scaling



Manual processes – slower times  
Static resource allocation



## Auto Scaling vs. Traditional Scaling

- **Traditional Scaling**
  - Provision servers to handle peak traffic, leading to underutilized resources during low traffic periods.
- **Auto Scaling**
  - Automatically adjusts capacity based on demand, ensuring cost-efficiency and maintaining performance.





# What is AWS Trusted Advisor?

A web service that inspects your AWS environment and provides real-time recommendations based on AWS best practices.

## Key Categories:

- Cost Optimization – Reducing AWS costs and improving resource usage.
- Performance – Improving the performance of AWS resources.
- Security – Identifying and addressing security vulnerabilities.
- Fault Tolerance – Enhancing availability and reliability.
- Service Limits – Alerts for approaching service usage limits.
- Dashboard Indicators:
  - Green Check: No issues detected.
  - Orange Triangle: Investigation recommended.
  - Red Circle: Action required.

Cost Optimization



0  9  0 

\$7,516.85

Potential monthly savings

Performance



3  7  0 

Security



2  4  11 




Fault Tolerance



0  15  5 

Service Limits



37  0  1 





**Cloud Trail**

## What is AWS CloudTrail?

- Records API calls in your AWS account, capturing detailed information such as the identity of the API caller, the time of the API call, and the source IP address.
- Functions as a “breadcrumb trail” of logs, allowing you to track and review all activities within your AWS resources.

### Key Features:

- Event History: Provides a complete history of user activities and API calls.
- Real-time Updates: Events are updated within 15 minutes after an API call.
- Filters: Filter events by time, date, user, resource type, etc., for better insights.



## **Example:**

If a new IAM user "Mary" is created but the owner doesn't know who made the change, CloudTrail will provide details such as:

**Who:** IAM user "John"

**When:** January 1, 2020, at 9:00 AM

**How:** Created through the AWS Management Console.

## **CloudTrail Insights:**

Automatically detects unusual API activities.

Example: If there's an unusual spike in Amazon EC2 instance launches, CloudTrail Insights will highlight this and allow further investigation.